

Heart Disease Prediction: An Applied Case Study

Rava Miro, Nazifi Endrit, Lucchina Luca, Del Grande Matteo
Data Science and AI
SUPSI

May 11, 2025

Abstract

The goal of the project is exploring and understanding heart disease datasets and predicting the presence/severity of the disease using machine learning techniques. We first performed a comprehensive review of available datasets, including the UCI Cleveland dataset and insights from existing research, highlighting challenges related to data quality and consistency. The primary analysis centered on the UCI Cleveland dataset, involving data cleaning, exploratory data analysis (EDA), feature engineering, and principal component analysis (PCA). We explored several classification models, including Naive Bayes, Random Forest (with various feature engineering), and a Multi-Layer Perceptron (MLP), were implemented and evaluated for binary classification (presence vs. absence of heart disease). The interpretability of the model was explored using SHAP in order to understand the important features affecting our predictions. Following our initial exploration, we analyzed the larger Sulianova cardiovascular disease dataset to compare model performance with the UCI Cleveland dataset.

Contents

1	Introduction: Project Premise	4
2	Data Understanding and Preparation (Cleveland Dataset)	5
2.1	Initial Data Loading and Cleaning	5
2.2	Feature Remapping	5
3	Exploratory Data Analysis (EDA) - Cleveland	5
3.1	Feature Distributions	5
3.2	Correlation Analysis	5
3.3	Outlier Analysis: Cholesterol ('chol')	5
4	Principal Component Analysis (PCA) - Cleveland	6
5	Modeling Strategy and Binary Classification (Cleveland)	6
5.1	Model Implementation and Evaluation (Binary Target)	6
5.1.1	Naive Bayes (GaussianNB)	6
5.1.2	Random Forest (Baseline)	6
5.1.3	Feature Engineering Iterations with Random Forest	6
5.1.4	Neural Network (MLP with PyTorch)	7
5.2	SHAP Analysis (Random Forest Baseline)	7
5.3	Model Comparison (Cleveland - Binary)	7
6	Extension: Analysis of Sulianova Dataset	8
6.1	Data Cleaning and Preprocessing (Sulianova)	8
6.2	EDA and Visualization (Sulianova)	8
6.3	Model Implementation and Evaluation (Sulianova)	8
7	Overall Conclusion and Future Work	9
8	Student Contributions	10

1 Introduction: Project Premise

Our project aims to predict the presence of heart disease using various patient features, following the guideline which specifies the target variable 'goal' as an integer value from 0 (no presence) to 4. Understanding that the foundation of any reliable model is high-quality, well-understood data is crucial as Data Scientists. Therefore, our initial step was to investigate the available datasets.

A common source for heart disease data is the UCI Machine Learning Repository, which contains several datasets like Cleveland, Hungary, Switzerland, and Long Beach VA. These originally had 76 attributes, but most analyses, including the widely cited ones, focus on a subset of 14 due to extensive missing data in the others.

However, the thesis by Brandon Simmons (2021) [1], has uncovered significant issues within these commonly used datasets. He investigated datasets from UCI, Kaggle, and Dataport, finding problems related to human errors in encoding, data duplication, and potentially misleading information stemming from these inconsistencies.

Key Findings from Simmons' Investigation

- **Kaggle vs. Cleveland Confusion:** A popular heart disease dataset found on Kaggle was closely related to the Cleveland dataset (303 records). Simons identified critical issues with the dataset as some missing data was imputed and, crucially, a reversed binary target variable encoding was found, leading to potentially misleading conclusions in many public analyses.
- **Statlog Dataset:** Another UCI dataset, Statlog (270 observations), was found by Simons to be a proper subset of the Cleveland dataset, highlighting risks of duplication if both were used together and duplicates were not handled properly.
- **Combined Datasets (e.g., ORIG/MISS):** Simmons examined combined datasets that merged multiple sources, including Cleveland and Statlog. He discovered high duplication rates and inconsistencies, particularly in the non-Cleveland/Statlog portions. These issues compromise data quality and the reliability of any models built on such datasets.
- **Recommendation:** Simmons concluded that the Cleveland dataset is the most reliable starting point, provided its few missing values are handled. He recommended recoding its 5-level target to binary (0 = Absence, 1 = Presence) and identified 'thalach', 'oldpeak', and 'cp' as key predictors.

Other Datasets Considered

We also reviewed other specific datasets:

- The **Fedesoriano - Heart Failure Prediction Dataset** on Kaggle [2] (918 observations, 11 features from 5 combined sources) presents similar duplication and quality concerns.
- The **Sulianova - Cardiovascular Disease Dataset** on Kaggle [3] (70,000 records) is much larger, with different features focused on general cardiovascular risk.

Our Approach

Following the advice of Brandon Simons, our primary analysis focused on the UCI Cleveland dataset. Our preprocessing pipeline included: handling missing values, ensuring correct data types, and deciding on the target variable representation (original 0-4 multi-class vs. binary 0/1). The Sulianova dataset was explored as an extension to the original analysis.

2 Data Understanding and Preparation (Cleveland Dataset)

We used the four heart disease datasets from UCI: Cleveland, Hungary, Switzerland, and VA. The Cleveland dataset was chosen for the primary analysis due to its relatively complete nature.

2.1 Initial Data Loading and Cleaning

The features that were available were the following: age, sex, cp, trestbps, chol, fbs, restecg, thalach, exang, oldpeak, slope, ca, thal, and num (target). Missing values, represented by '?', were converted to NaN. The Cleveland dataset initially had 303 rows. After identifying and removing rows with any missing values (6 rows in total, affecting 'ca' and 'thal'), the cleaned Cleveland dataset comprised 297 rows and 14 columns. Duplicate rows were checked across all datasets; the Cleveland dataset had no duplicates after initial loading, while others did, we did not explore these datasets, focused on the UCI Cleveland.

2.2 Feature Remapping

The 'thal' (Thallium stress test result) feature in the cleaned Cleveland dataset was remapped for clinical relevance and modeling suitability: normal (original value 3) was mapped to 0, reversible defect (original value 7) to 1, and fixed defect (original value 6) to 2.

3 Exploratory Data Analysis (EDA) - Cleveland

3.1 Feature Distributions

Distributions of all 13 features in the cleveland dataset were visualized using histograms for continuous features and bar plots for categorical ones. This helped understand the spread, central tendency, and potential imbalances for each feature.

3.2 Correlation Analysis

A correlation matrix heatmap was generated for all features in cleveland dataset.

- **Correlation with Target ('num'):** Features showing the strongest positive linear correlation with disease severity (num) were ca (number of major vessels, $r=0.52$), oldpeak (ST depression, $r=0.50$), and the remapped thal ($r=0.46$). thalach (max heart rate) showed a moderate negative correlation ($r=-0.42$). Features like fbs and chol had negligible linear correlation with num. These stronger correlations suggest that variables such as ca, oldpeak, and thal may be influential predictors in models that rely on linear relationships, though actual predictive importance should be confirmed through model-based feature importance metrics which will be explored later.
- **Inter-feature Correlation:** Strong correlations were noted between 'oldpeak' and 'slope' ($r=0.58$). 'thalach' was negatively correlated with 'slope', 'exang', and 'oldpeak'.

3.3 Outlier Analysis: Cholesterol ('chol')

The 'chol' feature was examined for outliers. While some high values were present (e.g., 1 patient with $\text{chol} > 500$ mg/dL, 4 patients > 400 mg/dL), a boxplot indicated these were few and considered clinically plausible extreme values rather than errors. Thus, they were retained.

4 Principal Component Analysis (PCA) - Cleveland

PCA was applied to the 13 scaled predictive features to visualize the data in lower dimensions.

- **Multi-class Target ('num' 0-4):** Both 2D and 3D PCA plots showed significant overlap between the 5 classes, with class 0 (no disease) being somewhat more distinct but no clear separation for classes 1-4. Graphs are available in the code and will be available in the presentation (we do not want to clutter the report with diagrams and images).
- **Binary Target (0 vs 1+):** A 3D PCA plot using a binary target (0 for no disease, 1 for any disease) showed a slightly clearer tendency for separation, but substantial overlap remained.
- **Cumulative Explained Variance:** The first three principal components 34.66%. To explain at least 90%, 10 components were needed; The last 3 components comprised the last 10% of the missing variance. This indicates that a few components do not capture the majority of the data's variance.

5 Modeling Strategy and Binary Classification (Cleveland)

Initial attempts with Random Forest on the multi-class target ('num' 0-4) showed poor performance, struggling on minority class predictions (severe disease stages). Consequently, the primary modeling strategy shifted to **binary classification**: predicting presence (1, if original 'num' > 0) versus absence (0, if original 'num' == 0) of heart disease. Data was split into train/test, stratified by the binary target. Numerical features were scaled using 'StandardScaler'.

5.1 Model Implementation and Evaluation (Binary Target)

A helper function was created to evaluate binary classifiers, reporting accuracy, precision, recall, F1-score (for the 'disease' class), AUC, and plotting confusion matrices and ROC curves for the test set. It also reported training accuracy.

5.1.1 Naive Bayes (GaussianNB)

A Gaussian Naive Bayes model was trained on the scaled baseline features.

- Test Accuracy: 90.00%
- AUC: 0.9375
- F1-score (Disease): 0.8889
- False Negatives (FN): 4

5.1.2 Random Forest (Baseline)

A Random Forest classifier (100 estimators, balanced class weight) was trained on the scaled baseline features.

- Test Accuracy: 86.67%
- AUC: 0.9408
- F1-score (Disease): 0.8525
- False Negatives (FN): 5

5.1.3 Feature Engineering Iterations with Random Forest

Several feature engineering techniques were explored, with a Random Forest model re-evaluated at each step:

1. **Ratio Features:** 'chol-per-age', 'tachycardia-index', 'bp-stress-ratio' were created. The RF model with these features showed slightly worse performance (Test Acc: 81.67%, AUC: 0.9401, FN: 8) compared to the baseline RF.
2. **Selected Features (7 features):** Based on statistical tests (ANOVA F-value, Chi-squared) and permutation importance (from RF with engineered features), a subset of 7 features ('thal', 'trestbps', 'ca', 'bp-stress-ratio', 'oldpeak', 'thalach', 'exang') was selected. We then trained a RF model which performed comparably to the baseline RF but with a slight dip in performance (Test Acc: 85.00%, AUC: 0.9158, FN: 6).
3. **Polynomial and Interaction Features (Degree 2):** Here we applied Polynomial features interaction of 2nd degree this was not beneficial as it led to worse RF performance (Test Acc: 81.67%, AUC: 0.9263, FN: 8) and signs of overfitting.
4. **Discretized (Binned) Features:** Binning 'oldpeak' and 'bp-stress-ratio' and using them with other categorical/engineered features significantly reduced RF performance (Test Acc: 75.00%, AUC: 0.8710, FN: 9).

These feature engineering experiments did not improve upon the baseline Random Forest or Naive Bayes models for the Cleveland dataset, increasing complexity without better generalization, even degrading the model's performance.

5.1.4 Neural Network (MLP with PyTorch)

A simple MLP (2 hidden layers with 64 and 32 neurons, ReLU, Batch Normalization, Dropout) was trained using PyTorch.

- **MLP with Baseline Features:** Test Accuracy: 83.33%, AUC: 0.9447, F1 (Disease): 0.8148, FN: 6.
- **MLP with Selected Features (7 features):** Test Accuracy: 83.33%, AUC: 0.9592, F1 (Disease): 0.7925, FN: 9. While achieving the highest AUC, its recall was lower, leading to more FNs, which in the medical field low FN values is very important.

5.2 SHAP Analysis (Random Forest Baseline)

SHAP was applied to the baseline Random Forest model to understand feature contributions.

- **Global Importance:** The global SHAP summary plot highlighted 'cp', 'thal', 'ca', 'oldpeak', and 'thalach' as highly impactful features in correct predictions.
- **Local Explanations:** SHAP plots for individual True Positive, False Negative, and False Positive cases illustrated how specific feature values pushed predictions for those instances.
- **Misclassifications:** Analysis of mean absolute SHAP values for FPs and FNs indicated which features were most influential in these incorrect predictions. For example, 'cp', 'ca', and 'thal' often had high importance in misclassified cases.

5.3 Model Comparison (Cleveland - Binary)

Table 1: Performance Summary for Binary Classification Models on Cleveland Dataset.

Model	Test Acc.	AUC	Precision (Dis.)	Recall (Dis.)	F1-score (Dis.)	FN
GaussianNB	0.900	0.938	0.923	0.857	0.889	4
RandomForest (Baseline)	0.867	0.941	0.885	0.821	0.852	5
RandomForest (Eng. Features)	0.817	0.940	0.870	0.714	0.784	8
RandomForest (Selected Features)	0.850	0.916	0.880	0.786	0.830	6
RandomForest (Poly & Interactions)	0.817	0.926	0.870	0.714	0.784	8
RandomForest (Binned Features)	0.750	0.871	0.760	0.679	0.717	9
PyTorch MLP (Baseline Feats)	0.833	0.945	0.846	0.786	0.815	6
PyTorch MLP (Selected Feats)	0.833	0.959	0.950	0.679	0.792	9

For the Cleveland dataset, GaussianNB provided the highest test accuracy and the lowest number of false negatives. The baseline Random Forest also performed well, particularly in AUC. The MLP with selected features achieved the highest AUC but had more false negatives. We can see that in the case of heart disease feature engineering and feature selection is not beneficial in terms of model performance.

6 Extension: Analysis of Sulianova Dataset

We explore the larger Sulanova Dataset, another cardiovascular disease dataset from Kaggle (including 70,000 records) which is composed of different features (e.g., age in days, height, weight, systolic/diastolic BP, lifestyle factors).

6.1 Data Cleaning and Preprocessing (Sulianova)

The 'id' column was dropped. Age was converted from days to years. Significant cleaning was performed for 'ap_hi' (systolic) and 'ap_lo' (diastolic) blood pressure, removing rows where 'ap_lo > ap_hi' or values were outside plausible physiological ranges (e.g., ap_hi 60-250, ap_lo 30-150). Outliers in height and weight were handled using percentile-based filtering (0.5th and 99.5th percentiles). We added a BMI feature computed based on the height, weight and age of the patient. The final processed dataset for Sulianova contained approximately 60,000-65,000 records (exact number depends on filtering effects described in notebook).

6.2 EDA and Visualization (Sulianova)

We started by plotting feature distributions to understand our data's behaviors, importantly seeing that the target variable 'cardio' was well-balanced.

- **PCA:** the first 2 principal components explained 40.2% of variance, adding the third increased it to 50.6%. As it can be seen in the Graphs (in the code) the classes were clustered quite noticeably. Cumulative variance analysis showed 7 components were needed for 80% variance, and 9 for 90%.
- **t-SNE:** It was applied to a sample of 5000 data points with a perplexity of 30, t-SNE revealed several distinct visual clusters in 2D and 3D, but these clusters still contained a significant mixture of both 'cardio' classes, indicating no clear non-linear separation.

6.3 Model Implementation and Evaluation (Sulianova)

We used a StandardScaler for numerical features and OneHotEncoded categorical features. Following that the models were trained on a 75%/25% train/test split.

Table 2: Performance Summary for Models on Sulianova Dataset.

Model	Test Acc.	AUC	Precision (Dis.)	Recall (Dis.)	F1-score (Dis.)	FN
GaussianNB (Sulianova)	0.697	0.769	0.744	0.591	0.658	3418
RandomForest (Sulianova)	0.713	0.770	0.715	0.697	0.706	2529
XGBoost (Sulianova)	0.734	0.800	0.753	0.689	0.719	2596
MLPClassifier (Sulianova)	0.735	0.805	0.764	0.681	0.712	2667

For the Sulianova dataset, MLPClassifier and XGBoost were the top performers, with MLPClassifier achieving the highest test accuracy, AUC, and precision. We noticed that the RandomForest showed signs of overfitting which was one of, if not the best performing model in the previous dataset.

7 Overall Conclusion and Future Work

This project successfully implemented and evaluated various machine learning models for heart disease prediction on two distinct datasets.

For the **UCI Cleveland dataset**, simpler models like **Gaussian Naive Bayes** (Test Acc: 90.0%, AUC: 0.938, FN: 4) and a **baseline Random Forest** (Test Acc: 86.7%, AUC: 0.941, FN: 5) demonstrated strong and balanced performance for binary classification. Iterative feature engineering attempts did not consistently yield improvements over these baseline models, sometimes degrading performance or increasing overfitting. SHAP analysis provided valuable insights into feature contributions seeing 'cp', 'thal', 'ca', 'oldpeak', and 'thalach'.

For the larger Sulianova dataset, more complex models like **MLPClassifier** (Test Acc: 0.735, AUC: 0.805) and **XGBoost** (Test Acc: 0.734, AUC: 0.800) performed best, showcasing their ability to handle larger, potentially more complex datasets. Visualization techniques like PCA and t-SNE highlighted the inherent class overlap in this dataset, outlining the difficulty of obtaining a high performing and robust model.

The key takeaways include the critical importance of thorough data understanding and preprocessing, the context-dependent success of feature engineering, and the realization that more complex models are not always superior, especially on smaller datasets. The choice of the "best" model also depends on the specific priorities (e.g., minimizing false negatives vs. maximizing overall accuracy).

Future Work could involve:

- Extensive hyperparameter tuning for all models, especially Random Forest, XGBoost, and MLP.
- Exploring other advanced classifiers (e.g., Support Vector Machines with different kernels, LightGBM).
- More sophisticated feature selection techniques (e.g., Recursive Feature Elimination).
- Investigating ensemble methods further by combining predictions from diverse models.
- Deeper analysis of misclassified instances to understand patterns and potential data issues.
- Addressing ethical implications and fairness in predictive healthcare models.

8 Student Contributions

- **Lucchina Luca**

- Dataset review and data quality issues (Cleveland, Kaggle, etc.)
- Research and summary of Simmons' findings
- Writing: *Introduction, Data Understanding*
- SHAP analysis (global and local)

- **Nazifi Endrit**

- Cleaning and preprocessing of the Cleveland dataset
- EDA: distributions, correlations, outliers
- PCA: analysis, visualizations, interpretation
- Support in model evaluation
- Feature engineering (ratios, polynomial features, selection)

- **Rava Miro**

- Experiments with Random Forest and MLP models
- Performance visualization (ROC curves, confusion matrix)
- Writing: *Modeling Strategy*
- Implementation: Naive Bayes models

- **Del Grande Matteo**

- Full analysis of the Sulianova dataset
- PCA and t-SNE on the extended dataset
- Writing: *Optional Extension, Conclusion and Future Work*
- Implementation of global models

- **Note**

- All group members contributed jointly to the implementation, training, and evaluation of the classification models (Naive Bayes, Random Forest, MLP, XGBoost) on both datasets.

References

- [1] Simmons, B. (2021). *Investigating Heart Disease Datasets and Building Predictive Models*. Master's Thesis, Elizabeth City State University. Available at: https://libres.uncg.edu/ir/ecsu/f/Brandon_Simmons_Thesis-Final.pdf
- [2] Fedesoriano. (n.d.). *Heart Failure Prediction Dataset*. Kaggle. Available at: <https://www.kaggle.com/datasets/fedesoriano/heart-failure-prediction>
- [3] Sulianova, S. (n.d.). *Cardiovascular Disease dataset*. Kaggle. Available at: <https://www.kaggle.com/datasets/sulianova/cardiovascular-disease-dataset>
- [4] Dua, D. and Graff, C. (2019). *UCI Machine Learning Repository* [<http://archive.ics.uci.edu/ml>]. Irvine, CA: University of California, School of Information and Computer Science. (General reference for UCI datasets).